

E-HANK: what the welfare numbers are, and what they are not

Working note for Boris and Francesca

August 2026

Boris asked whether the CEV numbers in the paper are subject to the well-known bias of first-order welfare evaluation (Kim and Kim, 2003; Benigno and Woodford, 2005). Short answer: the numbers *are* first-order approximations to welfare, so the concern is legitimate. They are not the stochastic case in which first-order evaluation is known to give spurious rankings, but the shock is large and the adoption margin is convex, so the size of the neglected second-order terms is an empirical question, which Section 4 settles with the nonlinear solver: levels are overstated by up to 10–15% at the headline shock, every policy ordering in the paper survives, and one sign claim (the crisis-only greening dividend) must be dropped. This note states precisely what the code computes, why it is the leading-order object, where it may fail, and what the check shows.

1. What `welfare.py` computes

Utilitarian welfare of discount-factor type i along a perfect-foresight transition indexed by shock size ε is

$$W_i(\varepsilon) = \sum_{t \geq 0} \beta_i^t \left[\int u(c_t(x; \varepsilon)) d\Lambda_{i,t}(x; \varepsilon) - v(n_t(\varepsilon)) \right], \quad v(n) = \varphi \frac{n^{1+1/\nu}}{1+1/\nu}, \quad (1)$$

where $x = (d, e, a)$ is the household state, $\Lambda_{i,t}$ the type- i distribution, and $v(n)$ the labour disutility implied by the union's wage-setting condition (uniform hours, so it is common to all households).¹

The code evaluates the transition part of (1) from the *linear* impulse response. `UTIL_i` is a het-output ($u(c)$ aggregated over the distribution); its linear IRF, returned by `solve_impulse_linear`, is by construction the first derivative of aggregate flow utility with respect to the shock,

$$d\text{UTIL}_{i,t} = \int u'(\bar{c}) \partial_\varepsilon c_t d\bar{\Lambda}_i + \int u(\bar{c}) \partial_\varepsilon \Lambda_{i,t}, \quad (2)$$

so that

$$dW_i = \underbrace{\frac{\text{UTIL}_i^{\text{pre}} - \text{UTIL}_i^{\text{base}} - dv^{\text{lvl}}}{1 - \beta_i}}_{\text{standing gap between steady states (exact)}} + \underbrace{\sum_t \beta_i^t (d\text{UTIL}_{i,t} - v'(\bar{n}) dn_t)}_{\text{crisis transition (first order in } \varepsilon)}, \quad (3)$$

and, with log utility, $\chi_i = \exp((1 - \beta_i) dW_i) - 1$, the permanent proportional consumption change that delivers dW_i . The mapping from dW_i to χ_i is exact; dW_i itself is $O(\varepsilon)$ with the $O(\varepsilon^2)$ terms dropped. The standing term (ETS steady state versus baseline steady state) comes from two nonlinear steady-state solves and is exact.

So the honest description of every CEV in the paper is: *the first-order Taylor expansion of welfare in the shock size, around a distorted steady state, along a deterministic transition*. What is dropped is

$$\frac{1}{2} \varepsilon^2 W_i'' = \frac{1}{2} \varepsilon^2 \sum_t \beta_i^t \left[\int (u''(\bar{c})(\partial_\varepsilon c_t)^2 + u'(\bar{c}) \partial_\varepsilon^2 c_t) d\bar{\Lambda}_i + 2 \int u'(\bar{c}) \partial_\varepsilon c_t \partial_\varepsilon \Lambda_{i,t} + \int u(\bar{c}) \partial_\varepsilon^2 \Lambda_{i,t} - v''(dn_t)^2 - v' \partial_\varepsilon^2 n_t \right]. \quad (4)$$

¹With `ghh_prefs=0` preferences are separable; $v(n)$ never enters the household's own problem and is added ex post from the aggregate n path.

2. Why this is nevertheless the leading-order object

The Kim–Kim “spurious welfare reversal” and the Benigno–Woodford LQ apparatus address the *stochastic* problem: expected welfare $\mathbb{E}W(\varepsilon)$ under recurring shocks. There $\mathbb{E}[\varepsilon W'] = 0$ by certainty equivalence, so the welfare effect of a policy is entirely second order, $\frac{1}{2}\sigma^2\mathbb{E}[W'']$, and a first-order solution gets W'' wrong because it omits $\partial_\varepsilon^2 c$ in (4). Rankings based on it can flip. Benigno and Woodford show that when the steady state is distorted, even the LQ objective requires a second-order approximation of the constraints.

Our exercise is different in one respect that matters: it is a single deterministic MIT transition. Then $\varepsilon W'_i \neq 0$ generically—the steady state is not Pareto-efficient (incomplete markets, markups, sticky wages), and the shock is a terms-of-trade income loss with first-order distributional incidence—so the first-order term is the dominant one and (3) captures it. The linear IRF is exactly the first-order perturbation solution (Boppart, Krusell and Mitman, 2018; Auclert, Bardóczy, Rognlie and Straub, 2021), so this is the standard object one obtains from a linearised HANK; it is not the object that Kim–Kim show to be spurious. For $\varepsilon \rightarrow 0$ the levels and rankings in the paper are correct to leading order.

3. Where it can fail: the shock is not small

Two facts in the paper qualify the “leading order” defence.

- The headline price shock doubles the world price ($\varepsilon = 1$). This is not a neighbourhood of the steady state.
- Appendix B documents that the adoption share is strongly convex in the shock: the nonlinear-to-linear ratio of the peak D^G response rises from 1.16 at $\varepsilon = 1/8$ to 2.04 at $\varepsilon = 3/4$, and the nonlinear solver does not converge at $\varepsilon = 1$. Aggregate output is essentially linear (ratio 1.00–1.01), which is why we have described the linear solution as reliable “for aggregate and consumption-based statistics”—a claim about allocations, not about welfare.

Orders of magnitude for the terms in (4). The curvature term $\frac{1}{2}u''(\bar{c})(dc)^2$ with consumption falling about 4–5% per quarter is roughly 2% of the first-order term—benign. The dangerous terms are those involving $\partial_\varepsilon^2 c$ and $\partial_\varepsilon^2 \Lambda$ through the adoption margin (switching costs paid, insulation gained, selection of who switches), which are exactly the ones the linear solution understates by up to a factor two.

The consequence is a distinction between the paper’s claims:

Comparison (price shock, CEV%)	gap	status at first order
transfer −0.82 vs cap −1.15	0.33 pp	probably robust; to be verified
flat transfer −0.32 vs Slutsky −0.82	0.50 pp	probably robust; to be verified
cap sign flip with supply elasticity	~ 0.7 pp	probably robust; to be verified
ex-ante ETS −1.852 vs laissez-faire −1.871	0.02 pp	inside the approximation error
greening dividend by persistence, +0.01 → −0.09 pp	≤ 0.1 pp	sign not established at first order

The first three rest on gaps that are themselves first order and large relative to any plausible second-order correction. The last two do not: “ex ante barely insures” is safe as a statement of equality, but “the greening dividend turns negative” (Section 6.4) is a sign claim resting on differences of a few basis points and should not be made from a first-order welfare calculation. Either we show it survives the nonlinear check or we restate it as “zero to first order”.

4. The nonlinear check

The sequence-space toolbox solves the nonlinear perfect-foresight transition (`solve_impulse_nonlinear`: Newton on the aggregate unknowns, household block re-solved exactly at every iteration). Along that path `UTIL.i` is the exact aggregate of $u(c)$ over the exact distribution, so (1) is evaluated without approximation. `run_nl_cev.py` does this for laissez-faire, cap, Slutsky transfer, flat transfer and the ex-ante ETS at $\varepsilon \in \{0.25, 0.5, 0.75\}$ (the headline $\varepsilon = 1$ does not converge; the cap does not converge at 0.75).

ε	Scenario	CEV% (linear)	CEV% (nonlinear)	NL/L	impat.	middle	patient
0.25	Laissez-faire	-0.471	-0.463	0.98	0.98	0.98	0.99
0.25	Ex-post cap	-0.288	-0.285	0.99	0.99	0.99	0.99
0.25	Slutsky transfer	-0.205	-0.198	0.97	0.95	0.97	0.99
0.25	Flat transfer	-0.079	-0.076	0.96	1.16	0.94	0.99
0.25	Ex-ante ETS	-0.434	-0.428	0.99	0.99	0.99	0.99
0.5	Laissez-faire	-0.940	-0.900	0.96	0.95	0.96	0.97
0.5	Ex-post cap	-0.575	-0.565	0.98	0.97	0.98	0.99
0.5	Slutsky transfer	-0.409	-0.377	0.92	0.88	0.92	0.96
0.5	Flat transfer	-0.159	-0.139	0.87	1.46	0.83	0.97
0.5	Ex-ante ETS	-0.909	-0.875	0.96	0.97	0.97	0.95
0.75	Laissez-faire	-1.407	-1.301	0.93	0.92	0.93	0.93
0.75	Slutsky transfer	-0.613	-0.524	0.86	0.79	0.85	0.91
0.75	Flat transfer	-0.238	-0.175	0.74	1.89	0.65	0.92
0.75	Ex-ante ETS	-1.382	-1.283	0.93	0.94	0.93	0.90

Four things follow.

1. *Levels.* The first-order CEV overstates the welfare loss, and the gap grows with the shock: 1–2% at $\varepsilon = 0.25$, 4–13% at 0.5, 7–26% at 0.75. The direction is the one the convexity of the adoption margin predicts: the true switching wave is larger (nonlinear/linear peak D^G of 1.3, 1.7, 2.0 at the three sizes) and cushions more. Extrapolating, the paper’s levels at $\varepsilon = 1$ are upper bounds on the losses, by something like 10–15% for laissez-faire and the cap and more for the transfers. They should be reported with that caveat.
2. *Ordering.* Flat > Slutsky > cap > ETS > laissez-faire at every size where all scenarios converge, linear and nonlinear alike. The first-order gaps barely move: Slutsky–cap +0.083 → +0.087 pp at 0.25 and +0.166 → +0.188 at 0.5; flat–Slutsky +0.125 → +0.122 and +0.250 → +0.238. The paper’s welfare rankings stand.
3. *Ex ante vs laissez-faire.* The gap is +0.037 → +0.034 pp at 0.25, +0.031 → +0.025 at 0.5, +0.025 → +0.018 at 0.75: a few basis points, equal to the standing gain of the ETS steady state (+0.043) minus a crisis-time dividend that is small and shrinking in ε . “Ex ante barely insures” is robust. The sign of the crisis-only dividend is not something a first-order calculation can establish, and Section 6.4 should be rephrased accordingly (“zero to first order, and if anything falling with the size and persistence of the shock”).
4. *Distribution.* The one place the linear solution errs in the other direction is the impatient type under the flat transfer: NL/L of 1.16, 1.46, 1.89. The linear solution understates the gain of high-MPC households when the transfer relaxes a binding liquidity constraint (asymmetry of u'). The distributional result of Section 5.5 (flat dominates energy-indexed targeting) is reinforced, not weakened, by the nonlinear check.

On precedents: we do not know of a HANK paper solved with the sequence-space Jacobian that derives a second-order welfare objective. Bhandari, Evans, Golosov and Sargent (2021)

do second-order (small-noise) welfare in HANK, outside SSJ, for optimal Ramsey policy under recurring shocks—the stochastic case. Dávila and Schaab (2025) provide the exact decomposition of heterogeneous-agent welfare changes and would be the natural frame if we go further than the ordering check.

References. Auclert, A., B. Bardóczy, M. Rognlie, L. Straub (2021), “Using the Sequence-Space Jacobian to Solve and Estimate Heterogeneous-Agent Models,” *Econometrica* 89(5), 2375–2408. Benigno, P., M. Woodford (2005), “Inflation Stabilization and Welfare: The Case of a Distorted Steady State,” *Journal of the European Economic Association* 3(6), 1185–1236. Benigno, P., M. Woodford (2012), “Linear-Quadratic Approximation of Optimal Policy Problems,” *Journal of Economic Theory* 147(1), 1–42. Bhandari, A., D. Evans, M. Golosov, T. Sargent (2021), “Inequality, Business Cycles, and Monetary-Fiscal Policy,” *Econometrica* 89(6), 2559–2599. Boppart, T., P. Krusell, K. Mitman (2018), “Exploiting MIT Shocks in Heterogeneous-Agent Economies,” *Journal of Economic Dynamics and Control* 89, 68–92. Dávila, E., A. Schaab (2025), “Welfare Assessments with Heterogeneous Individuals,” *Journal of Political Economy* 133(9), 2918–2961 (NBER WP 30571, 2022). Kim, J., S. H. Kim (2003), “Spurious Welfare Reversals in International Business Cycle Models,” *Journal of International Economics* 60(2), 471–500. Schmitt-Grohé, S., M. Uribe (2004), “Solving Dynamic General Equilibrium Models Using a Second-Order Approximation to the Policy Function,” *Journal of Economic Dynamics and Control* 28(4), 755–775.